XAI-Powered Fraud Detection

Ashish Seru
1BM21CS035

Tanisha Gotadke
1BM21CS229

Sneha Santhosh Bhat
1BM21CS213

Archit Mehrotra
1BM21CS031

Guide
Dr. Latha N. R.
Associate Professor

Department of Computer Science and Engineering, B.M.S. College of Engineering, Bangalore-560019

Introduction

Financial fraud is a growing concern in digital transactions, leading to significant financial losses and regulatory challenges. Traditional fraud detection systems, often rule-based or statistical, struggle with high false positives, evolving fraud tactics, and a lack of transparency in decision-making. This project proposes a Hybrid Explainable AI (XAI) Fraud Detection System that integrates machine learning, anomaly detection, and interpretability techniques to enhance fraud detection accuracy while ensuring regulatory compliance. By leveraging SHAP (SHapley Additive Explanations) the system provides clear insights into fraud detection decisions, making AI-driven processes more trustworthy and transparent. The proposed solution is designed to improve fraud detection efficiency, minimize operational costs, and increase stakeholder confidence in financial institutions. This project aims to build a highly accurate fraud detection system that minimizes false positives and negatives through advanced machine learning techniques. It incorporates anomaly detection using unsupervised learning to uncover suspicious transaction patterns and leverages Explainable AI (XAI), specifically SHAP, to provide clear justifications for fraud predictions. The system is designed to reduce financial losses and operational costs by improving detection efficiency, while also being scalable to handle large volumes of transactions. By offering transparency and interpretability, the project seeks to enhance user trust and confidence in automated fraud detection.

Proposed System

To overcome the shortcomings of traditional fraud detection systems, this project introduces a scalable, Explainable AI (XAI)-driven solution designed to enhance detection accuracy while ensuring transparency and compliance with financial regulations. Unlike static rule-based or opaque black-box models, our system adapts to evolving fraud patterns and clearly explains its decisions. It begins with comprehensive data preprocessing—cleaning and normalizing transactional data, engineering behavioral features (like transaction frequency, location patterns, and device usage), imputing missing values with SimpleImputer, scaling features using StandardScaler, and encoding class labels with LabelEncoder. Anomalies are flagged using Isolation Forest, and final fraud classification is performed using Random Forest. What sets this system apart is the use of SHAP (SHapley Additive Explanations), which provides granular, feature-level insights into each prediction, improving interpretability for analysts and regulators. An interactive Streamlit dashboard enables real-time monitoring and explanation of fraud decisions. The solution is not only accurate and interpretable but also deployable across cloud or on-premise infrastructure, making it a practical and trustworthy tool for modern financial environments.

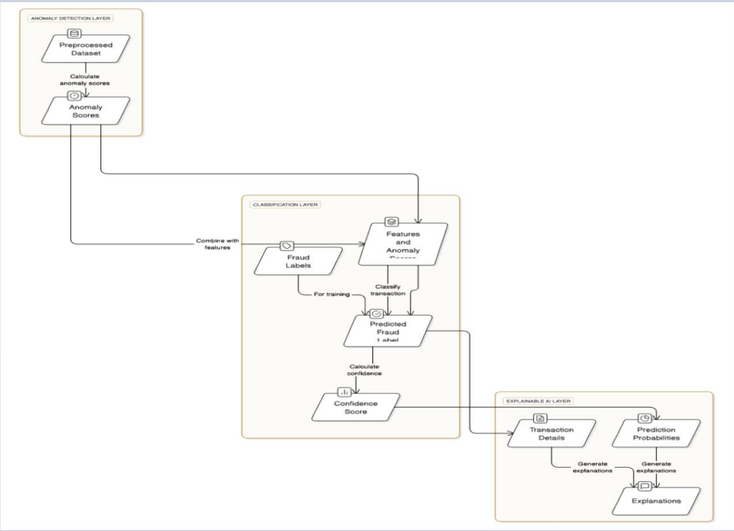

Fig 1. High level diagram of proposed system

Implementation details of modules

1. Data Ingestion Module:

This module is responsible for reading transaction data uploaded by users in CSV format. It utilizes the pandas library to efficiently load the data into the system. The module performs schema validation to ensure that all required columns are present and correctly formatted. User interaction is handled through Streamlit's file_uploader widget, providing a seamless interface for uploading files within the web application.

2. Preprocessing Module:

The preprocessing module cleans and prepares raw transaction data to be compatible with the machine learning models. Missing values in the dataset are handled using the SimpleImputer from sklearn.impute, which fills gaps using appropriate strategies such as mean, median, or most frequent values. Numeric features are standardized with StandardScaler from sklearn.preprocessing to normalize data ranges, improving model performance. Categorical variables, such as transaction channels or device identifiers, are encoded into numeric labels using LabelEncoder. To maintain consistency between training and prediction phases, these encoders are serialized and saved using the pickle library for later reuse.

3. Anomaly Detection Module:

This module detects unusual or suspicious transaction patterns by employing the Isolation Forest algorithm from sklearn.ensemble. The Isolation Forest isolates anomalies based on their attribute values, efficiently identifying transactions that deviate significantly from normal behavior. Anomaly scores are computed for each transaction, flagging those that are potentially fraudulent. This step can be optionally integrated into the classification pipeline to improve the overall detection accuracy.

4. Classification Module:

The classification module is tasked with categorizing each transaction as either fraudulent or legitimate. It uses the Random Forest classifier, an ensemble learning method known for its robustness and ability to handle heterogeneous data types. The model is trained on labeled historical transaction data and saved for deployment. During inference, the pre-trained model is loaded and used to predict the class of incoming transactions based on their features.

5. Explainability Module (XAI Layer):

To provide transparency and trust, this module implements explainability using SHAP (SHapley Additive exPlanations). SHAP values quantify the contribution of each feature to the model's prediction for an individual transaction. The module generates SHAP summary plots and force plots, which are integrated into the frontend interface to visually communicate how each feature influenced the fraud detection decision. This empowers analysts and stakeholders to understand and validate the system's outputs.

6. User Interface Module:

This module offers a user-friendly web interface built using Streamlit. It allows users to upload transaction files easily via the st.file_uploader() widget and preview data tables with st.dataframe(). The module also renders prediction results and SHAP-based explanation plots using st.pyplot() and shap.plots, enabling users to interact with and interpret the model's outputs intuitively. The interface is designed to require minimal technical knowledge and user effort for complete execution of the fraud detection pipeline.

7. Model & Encoder Management Module:

This module manages the lifecycle of machine learning models and encoders, ensuring they are properly saved and loaded across different sessions. It uses Python's pickle module to serialize pre-trained models and encoding objects into files. This process guarantees that the exact preprocessing steps and model parameters are consistently applied during inference, maintaining the integrity and reliability of predictions in production environments.

Data Flow of the System

The data flow in the fraud detection system begins with the user uploading transaction data files in CSV format through a web interface developed using Streamlit. Once uploaded, the Data Ingestion Module reads the file using the pandas library, validating the presence and correctness of essential columns and ensuring that the data adheres to the expected schema to prevent errors downstream. The validated raw data is then passed to the Preprocessing Module, where missing values are handled through imputation using SimpleImputer, numeric features are scaled uniformly with StandardScaler, and categorical features such as transaction channels and device identifiers are encoded into numeric labels via pre-trained LabelEncoder instances. These encoders are loaded from serialized files to maintain consistency with the training phase. The cleaned and transformed data proceeds to the Anomaly Detection Module, which applies the Isolation Forest algorithm to calculate anomaly scores for each transaction and flag unusual or suspicious patterns that may indicate fraud. Flagged anomalies, or optionally all transactions depending on system configuration, are forwarded to the Classification Module. Here, a Random Forest classifier, previously trained and serialized for deployment, analyzes each transaction's features and predicts its fraud status—fraudulent or legitimate. To ensure transparency and build stakeholder trust, the Explainability Module calculates SHAP (SHapley Additive exPlanations) values, quantifying the contribution of individual features to each prediction. These SHAP explanations, including summary and force plots, are integrated into the frontend dashboard to provide users with clear insights into why certain transactions were flagged. Finally, the Streamlit-based User Interface presents the complete results, allowing users to upload data, view prediction outcomes, and interactively explore the feature-level explanations with minimal effort. Throughout this entire workflow, all models and encoders are managed via persistent storage using Python's pickle module, enabling consistent and reproducible preprocessing and inference across different sessions and deployments.

Results

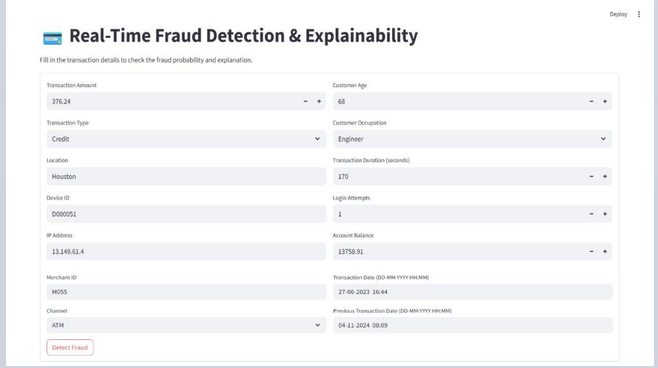

Fig 2. High level diagram of proposed system

Conclusion

The XAI-powered Fraud Detection System we developed is a significant step forward in tackling financial fraud. By integrating powerful machine learning models with explainability features like SHAP and LIME, the system offers not only high detection accuracy but also transparency, which is essential for building trust in AI-driven decisions.

Our system successfully minimized false positives, reduced financial losses, and helped analysts understand why certain transactions were flagged as fraudulent. This is a crucial advantage over traditional systems that often operate as "black boxes" without providing clear reasons for their decisions.

The solution we've built is scalable and can be integrated into existing financial systems, making it a valuable tool for preventing fraud while maintaining compliance with regulatory standards. By enhancing the interpretability of AI decisions, we've laid the groundwork for a safer, more secure financial ecosystem.

RESEARCH POSTER PRESENTATION DESIGN © 2012
www.PosterPresentations.com